

MACHINE LEARNING DEPLOYMENT USING IBM WATSON STUDIO

TEAM LEADER:

1.NAME:NIRUPAMA K

REG NO:211521243110

TEAM MEMBER:

2.NAME:NEHA SHRUTHI.U

REG NO:211521243109

3. NAME:ASWATHI SWARNA SREE

REG NO:211521243025

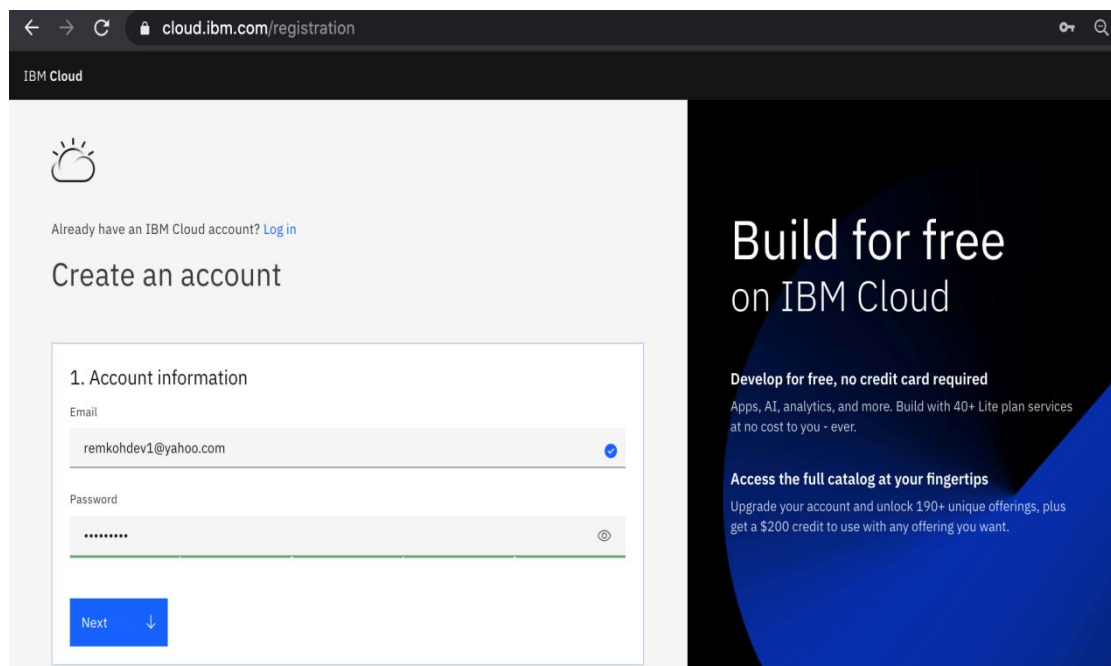
4. NAME:PRIYADHARSHINI N

REG NO: 211521243121

MACHINE LEARNING MODEL DEPLOYMENT USING IBM WATSON STUDIO

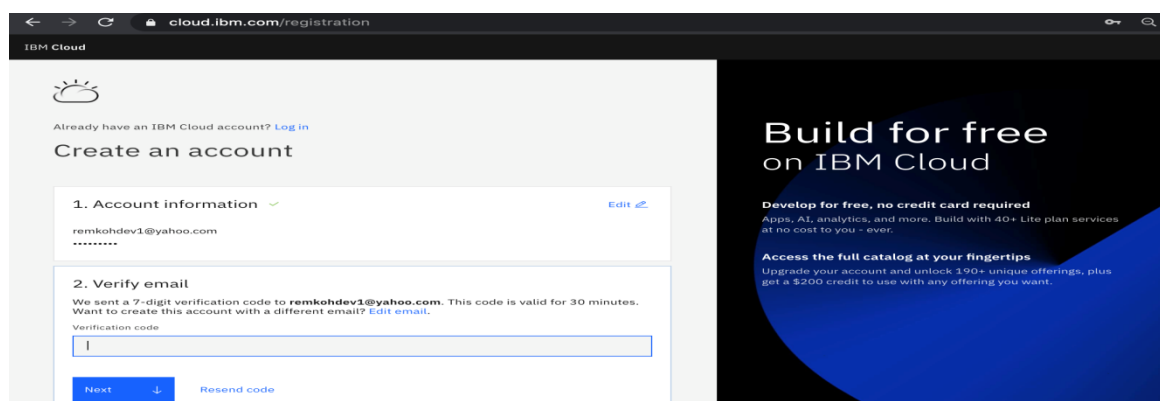
To create a new account, follow the steps below,

1. You need an IBM Cloud account to access your cluster,
2. If you do not have an IBM Cloud account yet, go to <https://cloud.ibm.com/registration> to register,
3. In the Create an account window, enter your email and password,



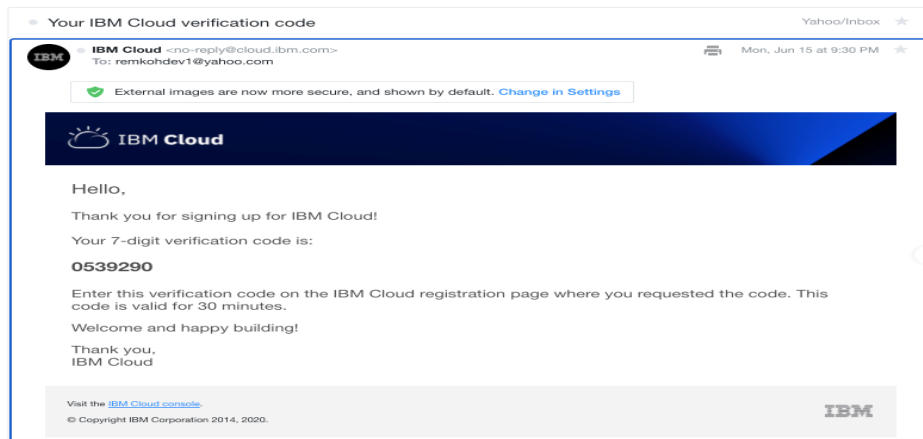
The screenshot shows the IBM Cloud registration page at cloud.ibm.com/registration. The page has a dark header with the IBM Cloud logo. Below the header, there's a light gray section with the IBM Cloud logo and the text "Already have an IBM Cloud account? [Log in](#)". The main heading is "Create an account". Below this, there's a form titled "1. Account information". It has two input fields: "Email" with the value "remkohdev1@yahoo.com" and a blue checkmark icon, and "Password" with a masked password "*****" and a blue eye icon. At the bottom of the form is a blue button labeled "Next" with a downward arrow.

4. The Verify email section will inform you that a verification code was sent to your email,



The screenshot shows the IBM Cloud registration page at cloud.ibm.com/registration. The page has a dark header with the IBM Cloud logo. Below the header, there's a light gray section with the IBM Cloud logo and the text "Already have an IBM Cloud account? [Log in](#)". The main heading is "Create an account". Below this, there's a form titled "1. Account information" with a green checkmark and an "Edit" link. It has two input fields: "Email" with the value "remkohdev1@yahoo.com" and "Password" with a masked password "*****". Below this, there's a section titled "2. Verify email". It contains the text "We sent a 7-digit verification code to remkohdev1@yahoo.com. This code is valid for 30 minutes. Want to create this account with a different email? [Edit email](#)." Below this text is a text input field for the "Verification code". At the bottom of the form are two buttons: a blue button labeled "Next" with a downward arrow, and a blue link labeled "Resend code".

5. Switch to your email provider to retrieve the verification code,



6. Enter the verification code in the verify email section, and click Next,

A screenshot of the IBM Cloud registration page in a web browser. The URL is cloud.ibm.com/registration. The page has a dark header with the IBM Cloud logo. On the left, there's a sidebar with the IBM Cloud logo and links for 'Log in' and 'Create an account'. The main content area has two steps: '1. Account information' (completed) and '2. Verify email'. In the 'Verify email' section, it says 'We sent a 7-digit verification code to remkohdev1@yahoo.com. This code is valid for 30 minutes. Want to create this account with a different email? Edit email.' Below this is a 'Verification code' input field containing '0539290'. At the bottom of the form are 'Next' and 'Resend code' buttons. On the right, there's a large blue and black graphic with the text 'Build for free on IBM Cloud' and 'Develop for free, no credit card required'.

7. Enter your first name, last name and country in the Personal information section and click Next,

8. Click Create account,

9. Your account is being created,

10. Review the IBM Privacy Statement,

About your IBMid Account Privacy

Updates to the [IBM Privacy Statement](#) since this notice was originally published provide additional information about how your personal information is processed by IBM.

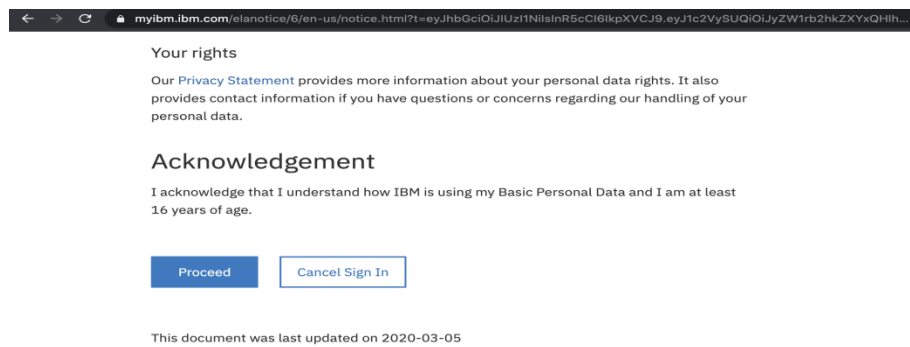
+ Changes since the previous version of this notice

- + What data does IBM collect?

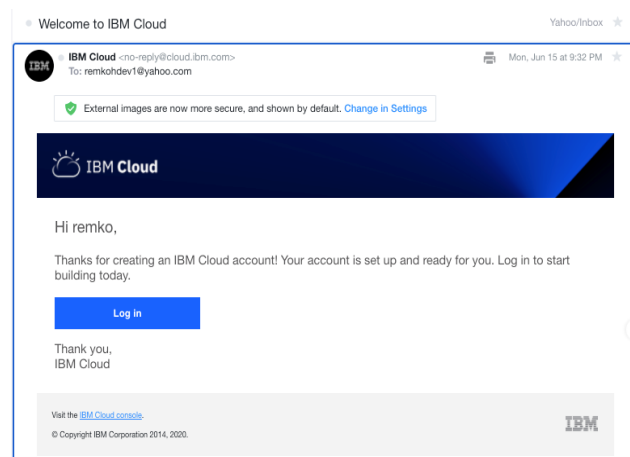
+ Why IBM needs your data

+ How your data was obtained

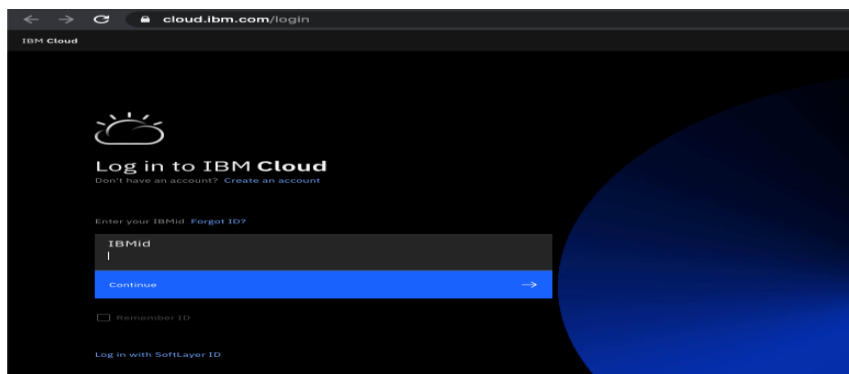
11. Click Proceed to acknowledge the privacy statement,




12. Switch to your email provider to review the Welcome to IBM Cloud email, and click the Login link,



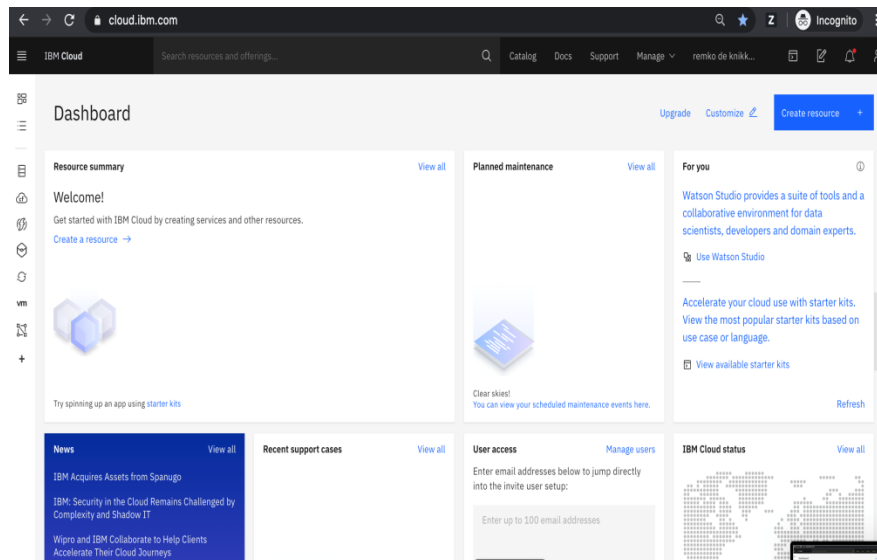
13. Enter your IBM Id to login,



14. Enter your password to login,



15. The IBM Cloud dashboard page should load,



16. You have successfully registered a new IBM Cloud account.

Churn Prediction

A Machine Learning Model That Can Predict Customers Who Will Leave The Company
The aim is to predict whether a bank's customers leave the bank or not. If the Client has closed his/her bank account, he/she has left.

Dataset:

- **RowNumber:** corresponds to the record (row) number and has no effect on the output.
- **CustomerId:** contains random values and has no effect on customer leaving the bank.
- **Surname:** the surname of a customer has no impact on their decision to leave the bank.
- **CreditScore:** can have an effect on customer churn, since a customer with a higher credit score is less likely to leave the bank.
- **Geography:** a customer's location can affect their decision to leave the bank.
- **Gender:** it's interesting to explore whether gender plays a role in a customer leaving the bank.
- **Age:** this is certainly relevant, since older customers are less likely to leave their bank than younger ones.

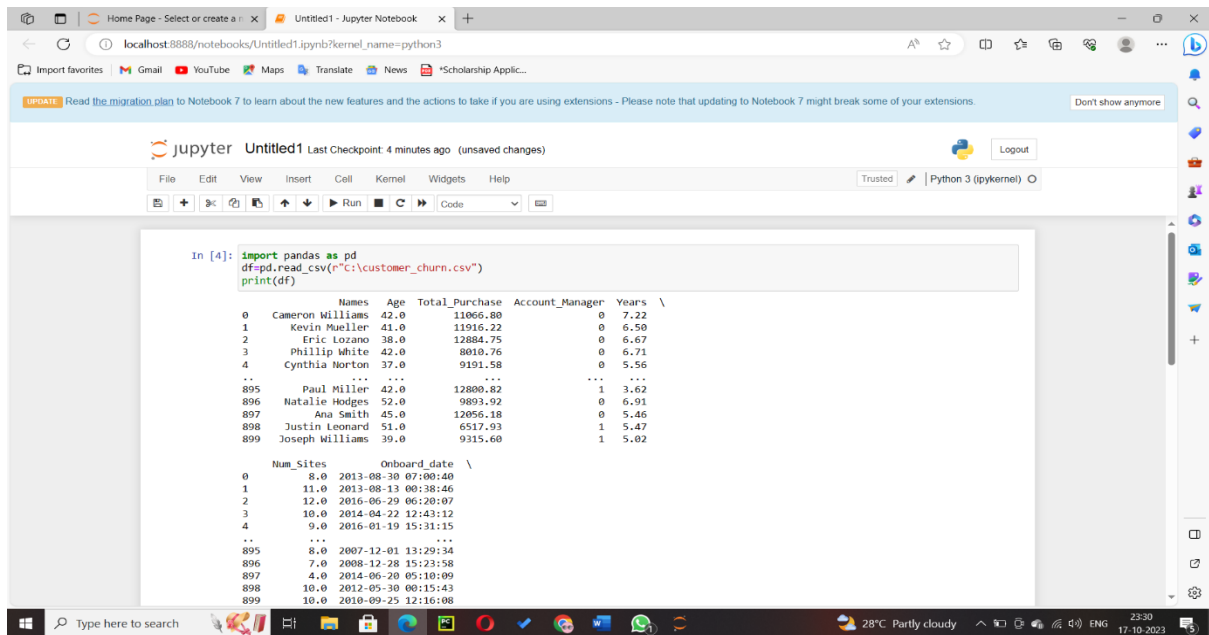
- **Tenure:** refers to the number of years that the customer has been a client of the bank. Normally, older clients are more loyal and less likely to leave a bank.
- **Balance:** also a very good indicator of customer churn, as people with a higher balance in their accounts are less likely to leave the bank compared to those with lower balances.
- **NumOfProducts:** refers to the number of products that a customer has purchased through the bank.
- **HasCrCard:** denotes whether or not a customer has a credit card. This column is also relevant, since people with a credit card are less likely to leave the bank.
- **IsActiveMember:** active customers are less likely to leave the bank.
- **EstimatedSalary:** as with balance, people with lower salaries are more likely to leave the bank compared to those with higher salaries.
- **Exited:** whether or not the customer left the bank. (0=No,1=Yes)

Code:

```
import pandas as pd

df=pd.read_csv(r"c:\customer_churn.csv")

print(df)
```



Exploratory Data Analysis:

The first thing we have to do in Exploratory Data Analysis is checked if there are null values in the dataset.

`df.isnull().head()`

	customer_id	credit_score	country	gender	age	tenure	balance	products_number	credit_card	active_member	estimated_salary	churn
0	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False

#This is used to check any null value.

`df.isnull().sum()`

```
customer_id      0
credit_score     0
country          0
gender           0
age              0
tenure           0
balance          0
products_number  0
credit_card      0
active_member    0
estimated_salary 0
churn            0
dtype: int64
```


#Checking Data types

df.dtypes

```
customer_id      int64
credit_score      int64
country          object
gender           object
age             int64
tenure           int64
balance          float64
products_number  int64
credit_card      int64
active_member    int64
estimated_salary float64
churn            int64
dtype: object
```

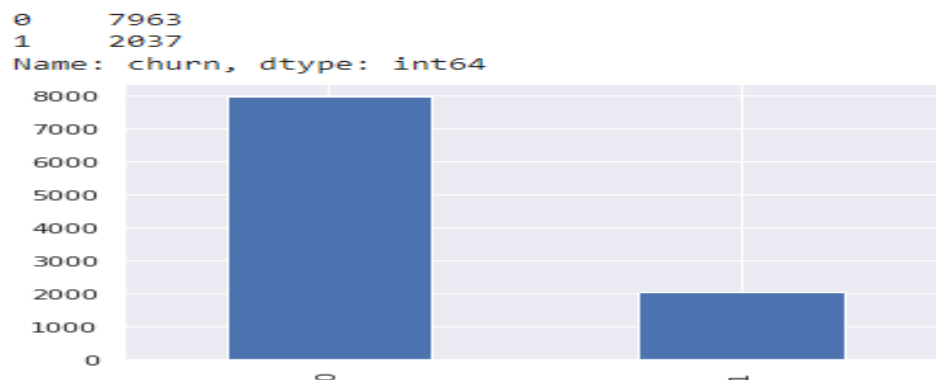
#Counting 1 and 0 Value in Churn column

```
color_wheel = {1: "#0392cf", 2: "#7bc043"}
```

```
colors = df["churn"].map(lambda x: color_wheel.get(x + 1))
```

```
print(df.churn.value_counts())
```

```
p=df.churn.value_counts().plot(kind="bar")
```



#Change value in country column

```
df['country'] = df['country'].replace(['Germany'],'0')
```

```
df['country'] = df['country'].replace(['France'],'1')
```

```
df['country'] = df['country'].replace(['Spain'],'2')
```

#Change value in gender column

```
df['gender'] = df['gender'].replace(['Female'],'0')
```

```
df['gender'] = df['gender'].replace(['Male'],'1')
```

```
df.head()
```

	customer_id	credit_score	country	gender	age	tenure	balance	products_number	credit_card	active_member	estimated_salary	churn
0	15634602	619	1	0	42	2	0.00	1	1	1	101348.88	1
1	15647311	608	2	0	41	1	83807.86	1	0	1	112542.58	0
2	15619304	502	1	0	42	8	159660.80	3	1	0	113931.57	1
3	15701354	699	1	0	39	1	0.00	2	0	0	93826.63	0
4	15737888	850	2	0	43	2	125510.82	1	1	1	79084.10	0

```
#convert object data types column to integer
```

```
df['country'] = pd.to_numeric(df['country'])
```

```
df['gender'] = pd.to_numeric(df['gender'])
```

```
df.dtypes
```

```
customer_id      int64
credit_score      int64
country          int64
gender           int64
age              int64
tenure           int64
balance          float64
products_number  int64
credit_card      int64
active_member    int64
estimated_salary float64
churn            int64
dtype: object
```

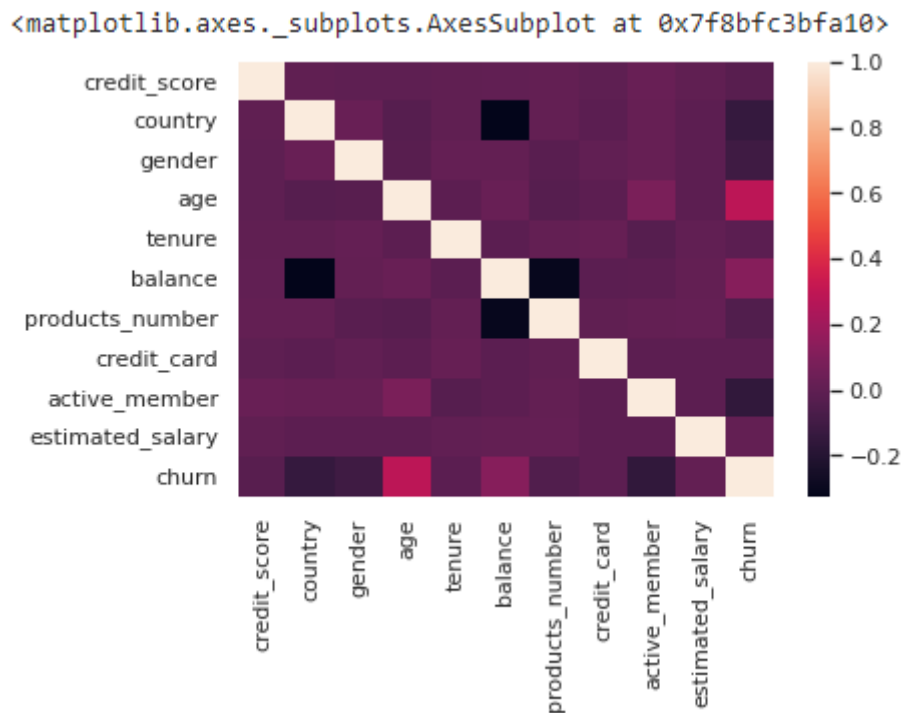
```
#Remove customer_id column
```

```
df2 = df.drop('customer_id', axis=1)
```

```
df2.head()
```

	credit_score	country	gender	age	tenure	balance	products_number	credit_card	active_member	estimated_salary	churn
0	619	1	0	42	2	0.00	1	1	1	101348.88	1
1	608	2	0	41	1	83807.86	1	0	1	112542.58	0
2	502	1	0	42	8	159660.80	3	1	0	113931.57	1
3	699	1	0	39	1	0.00	2	0	0	93826.63	0
4	850	2	0	43	2	125510.82	1	1	1	79084.10	0

```
sns.heatmap(df2.corr(), fmt='.2g')
```



Build Machine Learning Model:

```
X = df2.drop('churn', axis=1)
```

```
y = df2['churn']
```

```
#test size 20% and train size 80%
```

```
from sklearn.model_selection import train_test_split, cross_val_score,
```

```
cross_val_predict
```

```
from sklearn.metrics import accuracy_score
```

```
X_train, X_test, y_train, y_test = train_test_split(X,y,
```

```
test_size=0.2,random_state=7)
```

Decision Tree:

```
from sklearn.tree import DecisionTreeClassifier
```

```
dtree = DecisionTreeClassifier()
```

```
dtree.fit(X_train, y_train)
```

```
y_pred = dtree.predict(X_test)
```

```
print("Accuracy Score :", accuracy_score(y_test, y_pred)*100, "%")
```

Accuracy Score : 79.0 %

Random Forest:

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(X_train, y_train)
y_pred = rfc.predict(X_test)
print("Accuracy Score :", accuracy_score(y_test, y_pred)*100, "%")
```

Accuracy Score : 86.85000000000001 %

Support Vector Machine:

```
from sklearn import svm
svm = svm.SVC()
svm.fit(X_train, y_train)
y_pred = svm.predict(X_test)
print("Accuracy Score :", accuracy_score(y_test, y_pred)*100, "%")
```

Accuracy Score : 79.45 %

XGBoost:

```
from xgboost import XGBClassifier
xgb_model = XGBClassifier()
xgb_model.fit(X_train, y_train)
y_pred = xgb_model.predict(X_test)
print("Accuracy Score :", accuracy_score(y_test, y_pred)*100, "%")
```

Accuracy Score : 86.45 %

Visualize Random Forest and XGBoost Algorithm because Random Forest and XGBoost Algorithm have the Best Accuracy:


```
#importing classification report and confusion matrix from sklearn  
from sklearn.metrics import classification_report, confusion_matrix
```

Random Forest:

```
y_pred = rfc.predict(X_test)  
print("Classification report - n", classification_report(y_test,y_pred))
```

```
Classification report -  
              precision    recall  f1-score   support  
  
     0           0.88       0.97       0.92       1589  
     1           0.80       0.48       0.60        411  
  
 accuracy              0.87       2000  
 macro avg           0.84       0.72       0.76       2000  
 weighted avg        0.86       0.87       0.86       2000
```


```
cm = confusion_matrix(y_test, y_pred)  
plt.figure(figsize=(5,5))  
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap =  
'Blues')  
plt.ylabel('Actual label')  
plt.xlabel('Predicted label')  
all_sample_title = 'Accuracy Score: {0}'.format(rfc.score(X_test, y_test))  
plt.title(all_sample_title, size = 15)
```



```

from sklearn.metrics import roc_curve, roc_auc_score
y_pred_proba = rfc.predict_proba(X_test)[:][:,1]
df_actual_predicted = pd.concat([pd.DataFrame(np.array(y_test),
columns=['y_actual']), pd.DataFrame(y_pred_proba,
columns=['y_pred_proba'])], axis=1)
df_actual_predicted.index = y_test.index
fpr, tpr, tr = roc_curve(df_actual_predicted['y_actual'],
df_actual_predicted['y_pred_proba'])
auc = roc_auc_score(df_actual_predicted['y_actual'],
df_actual_predicted['y_pred_proba'])
plt.plot(fpr, tpr, label='AUC = %0.4f' % auc)
plt.plot(fpr, fpr, linestyle = '--', color='k')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve', size = 15)
plt.legend()

```



XGBoost:

```
y_pred = xgb_model.predict(X_test)
```

```
print("Classification report - n", classification_report(y_test,y_pred))
```

```
Classification report -
              precision    recall  f1-score   support

     0       0.87        0.97        0.92        1589
     1       0.80        0.46        0.58         411

 accuracy          0.86          2000
 macro avg         0.84          0.71          0.75          2000
 weighted avg      0.86          0.86          0.85          2000
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(5,5))
```


```
sns.heatmap(data=cm,linewidths=.5, annot=True,square = True, cmap =
'Blues')
```

```
plt.ylabel('Actual label')
```

```
plt.xlabel('Predicted label')
```

```
all_sample_title = 'Accuracy Score: {0}'.format(xgb_model.score(X_test,
y_test))
```


```
plt.title(all_sample_title, size = 15)
```



```

from sklearn.metrics import roc_curve, roc_auc_score
y_pred_proba = xgb_model.predict_proba(X_test)[:][:,1]
df_actual_predicted = pd.concat([pd.DataFrame(np.array(y_test),
columns=['y_actual']), pd.DataFrame(y_pred_proba,
columns=['y_pred_proba'])], axis=1)
df_actual_predicted.index = y_test.index
fpr, tpr, tr = roc_curve(df_actual_predicted['y_actual'],
df_actual_predicted['y_pred_proba'])
auc = roc_auc_score(df_actual_predicted['y_actual'],
df_actual_predicted['y_pred_proba'])
plt.plot(fpr, tpr, label='AUC = %0.4f' % auc)
plt.plot(fpr, fpr, linestyle = '--', color='k')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve', size = 15)
plt.legend()

```

```
# Descriptive statistics of the data set
df.describe([0.05,0.25,0.50,0.75,0.90,0.95,0.99])
```

Out[6]:

	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402769
min	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000	0.000000
5%	1.557882e+07	489.000000	25.000000	1.000000	0.000000	1.000000	0.00000	0.000000	9851.818500	0.000000
25%	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000	0.000000

	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
50 %	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.000000	1.000000	100193.915000	0.000000
75 %	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.000000	1.000000	149388.247500	0.000000
90 %	1.579083e+07	778.000000	53.000000	9.000000	149244.792000	2.000000	1.000000	1.000000	179674.704000	1.000000
95 %	1.580303e+07	812.000000	60.000000	9.000000	162711.669000	2.000000	1.000000	1.000000	190155.375500	1.000000
99 %	1.581311e+07	850.000000	72.000000	10.000000	185967.985400	3.000000	1.000000	1.000000	198069.734500	1.000000
max	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000				

categorical Variables

```
categorical_variables = [col for col in df.columns if col in "0"
                        or df[col].nunique() <=11
                        and col not in "Exited"]
```

categorical_variables

Out[7]:

```
['Geography',
 'Gender',
 'Tenure',
 'NumOfProducts',
 'HasCrCard',
 'IsActiveMember']
```

In [8]:

Numeric Variables

```
numeric_variables = [col for col in df.columns if df[col].dtype != "object"
                    and df[col].nunique() >11
                    and col not in "CustomerId"]
```

numeric_variables

Out[8]:

```
['CreditScore', 'Age', 'Balance', 'EstimatedSalary']
```

Exited (Dependent Variable)

In [9]:

```
# Frequency of classes of dependent variable  
df["Exited"].value_counts()
```

Out[9]:

```
0    7963  
1    2037  
Name: Exited, dtype: int64
```

In [10]:

```
# Customers Leaving the bank  
churn = df.loc[df["Exited"]==1]
```

In [11]:

```
# Customers who did not leave the bank  
not_churn = df.loc[df["Exited"]==0]
```

Categorical Variables

Tenure

In [12]:

```
# Frequency of not_churn group according to Tenure  
not_churn["Tenure"].value_counts().sort_values()
```

Out[12]:

```
0    318  
10   389  
6    771  
9    771  
4    786  
3    796  
5    803  
1    803  
8    828  
2    847  
7    851  
Name: Tenure, dtype: int64
```

In [13]:

```
# Frequency of churn group according to Tenure  
churn["Tenure"].value_counts().sort_values()
```

Out[13]:

```
0     95  
10    101  
7     177  
6     196  
8     197  
2     201  
4     203  
5     209  
9     213
```

```
3      213
1      232
Name: Tenure, dtype: int64
```

NumOfProducts

In [14]:

```
# Frequency of not_churn group according to NumOfProducts
not_churn["NumOfProducts"].value_counts().sort_values()
```

Out[14]:

```
3      46
1     3675
2     4242
Name: NumOfProducts, dtype: int64
```

In [15]:

```
# Frequency of churn group according to NumOfProducts
churn["NumOfProducts"].value_counts().sort_values()
```

Out[15]:

```
4      60
3     220
2     348
1    1409
Name: NumOfProducts, dtype: int64
```

HasCrCard

In [16]:

```
# examining the HasCrCard of the not_churn group
not_churn["HasCrCard"].value_counts()
```

Out[16]:

```
1     5631
0     2332
Name: HasCrCard, dtype: int64
```

In [17]:

```
# examining the HasCrCard of the churn group
churn["HasCrCard"].value_counts()
```

Out[17]:

```
1     1424
0      613
Name: HasCrCard, dtype: int64
```

IsActiveMember

In [18]:

```
# examining the IsActiveMember of the not_churn group
not_churn["IsActiveMember"].value_counts()
```

Out[18]:

```
1     4416
0     3547
Name: IsActiveMember, dtype: int64
```

In [19]:

```
# examining the IsActiveMember of the churn group
churn["IsActiveMember"].value_counts()
```

```
Out[19]:
```

```
0    1302
```

```
1     735
```

```
Name: IsActiveMember, dtype: int64
```

Geography

```
In [20]:
```

```
# Frequency of not_churn group according to Geography
```

```
not_churn.Geography.value_counts().sort_values()
```

```
Out[20]:
```

```
Germany    1695
```

```
Spain      2064
```

```
France     4204
```

```
Name: Geography, dtype: int64
```

```
In [21]:
```

```
# Frequency of churn group according to Geography
```

```
churn.Geography.value_counts().sort_values()
```

#Reference dataset: <https://www.kaggle.com/datasets/gauravtopre/bank-custom-er-churn-dataset>

Conclusion:

Developing a Machine Learning Model is a complex process, but it is essential for building and deploying successful machine-learning applications. By following the steps outlined in this blog, you can increase your chances of success.